

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline

from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score

data= pd.read_csv(r"C:\Users\santo\OneDrive\Desktop\Data science\Dec2025\Machine Learning\car-mpg.csv")
```

```
In [3]: data.head()
```

```
Out[3]:   mpg cyl disp hp wt acc yr origin car_type          car_name
0    18.0   8  307.0 130  3504 12.0 70      1       0  chevrolet chevelle malibu
1    15.0   8  350.0 165  3693 11.5 70      1       0        buick skylark 320
2    18.0   8  318.0 150  3436 11.0 70      1       0  plymouth satellite
3    16.0   8  304.0 150  3433 12.0 70      1       0        amc rebel sst
4    17.0   8  302.0 140  3449 10.5 70      1       0        ford torino
```

```
In [4]: data.head(20)
```

Out[4]:

	mpg	cyl	disp	hp	wt	acc	yr	origin	car_type	car_name
0	18.0	8	307.0	130	3504	12.0	70	1	0	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	0	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	0	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	0	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	0	ford torino
5	15.0	8	429.0	198	4341	10.0	70	1	0	ford galaxie 500
6	14.0	8	454.0	220	4354	9.0	70	1	0	chevrolet impala
7	14.0	8	440.0	215	4312	8.5	70	1	0	plymouth fury iii
8	14.0	8	455.0	225	4425	10.0	70	1	0	pontiac catalina
9	15.0	8	390.0	190	3850	8.5	70	1	0	amc ambassador dpl
10	15.0	8	383.0	170	3563	10.0	70	1	0	dodge challenger se
11	14.0	8	340.0	160	3609	8.0	70	1	0	plymouth 'cuda 340
12	15.0	8	400.0	150	3761	9.5	70	1	0	chevrolet monte carlo
13	14.0	8	455.0	225	3086	10.0	70	1	0	buick estate wagon (sw)
14	24.0	4	113.0	95	2372	15.0	70	3	1	toyota corona mark ii
15	22.0	6	198.0	95	2833	15.5	70	1	0	plymouth duster
16	18.0	6	199.0	97	2774	15.5	70	1	0	amc hornet
17	21.0	6	200.0	85	2587	16.0	70	1	0	ford maverick
18	27.0	4	97.0	88	2130	14.5	70	3	1	datsun pl510
19	26.0	4	97.0	46	1835	20.5	70	2	1	volkswagen 1131 deluxe sedan

In [12]:

```
import numpy as np
import pandas as pd
```

```
# clean column names
data.columns = data.columns.str.strip().str.lower()

# drop car_name safely
data.drop(columns=['car_name'], errors='ignore', inplace=True)

# encode origin only if present
if 'origin' in data.columns:
    data['origin'] = data['origin'].map({
        1: 'america',
        2: 'europe',
        3: 'asia'
    })
    data = pd.get_dummies(data, columns=['origin'], dtype=int)

# replace missing symbols
data.replace('?', np.nan, inplace=True)
```

```
In [14]: import numpy as np
import pandas as pd

data = data.apply(pd.to_numeric, errors='ignore')

numeric_cols = data.select_dtypes(include=[np.number]).columns

data[numeric_cols] = data[numeric_cols].apply(
    lambda x: x.fillna(x.median()))
)
```

```
C:\Users\santo\AppData\Local\Temp\ipykernel_4516\3584547213.py:5: FutureWarning: errors='ignore' is deprecated and will raise in a future version. Use to_numeric without passing `errors` and catch exceptions explicitly instead
  data = data.apply(pd.to_numeric, errors='ignore')
```

```
In [15]: data.head()
```

Out[15]:

	mpg	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe
0	18.0	8	307.0	130.0	3504	12.0	70	0	1	0	0
1	15.0	8	350.0	165.0	3693	11.5	70	0	1	0	0
2	18.0	8	318.0	150.0	3436	11.0	70	0	1	0	0
3	16.0	8	304.0	150.0	3433	12.0	70	0	1	0	0
4	17.0	8	302.0	140.0	3449	10.5	70	0	1	0	0

In [16]:

data

Out[16]:

	mpg	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe
0	18.0	8	307.0	130.0	3504	12.0	70	0	1	0	0
1	15.0	8	350.0	165.0	3693	11.5	70	0	1	0	0
2	18.0	8	318.0	150.0	3436	11.0	70	0	1	0	0
3	16.0	8	304.0	150.0	3433	12.0	70	0	1	0	0
4	17.0	8	302.0	140.0	3449	10.5	70	0	1	0	0
...
393	27.0	4	140.0	86.0	2790	15.6	82	1	1	0	0
394	44.0	4	97.0	52.0	2130	24.6	82	1	0	0	1
395	32.0	4	135.0	84.0	2295	11.6	82	1	1	0	0
396	28.0	4	120.0	79.0	2625	18.6	82	1	1	0	0
397	31.0	4	119.0	82.0	2720	19.4	82	1	1	0	0

398 rows × 11 columns

2. Model building

Here we would like to scale the data as the columns are varied which would result in 1 column dominating the others.

First we divide the data into independent (X) and dependent data (y) then we scale it.

Tip!:

*The reason we don't scale the entire data before and then divide it into train(X) & test(y) is because once you scale the data, the type(data_s) would be numpy.ndarray. It's impossible to divide this data when it's an array. *

Hence we divide type(data) pandas.DataFrame, then proceed to scaling it.

```
In [17]: x = data.drop(['mpg'], axis = 1)
y = data[['mpg']]
```

```
In [21]: x_s = preprocessing.scale(x)
x_s = pd.DataFrame(x_s, columns = x.columns )

y_s = preprocessing.scale(y)
y_s = pd.DataFrame(y_s, columns = y.columns)
```

```
In [24]: x_s
```

Out[24]:

	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe
0	1.498191	1.090604	0.673118	0.630870	-1.295498	-1.627426	-1.062235	0.773559	-0.497643	-0.461968
1	1.498191	1.503514	1.589958	0.854333	-1.477038	-1.627426	-1.062235	0.773559	-0.497643	-0.461968
2	1.498191	1.196232	1.197027	0.550470	-1.658577	-1.627426	-1.062235	0.773559	-0.497643	-0.461968
3	1.498191	1.061796	1.197027	0.546923	-1.295498	-1.627426	-1.062235	0.773559	-0.497643	-0.461968
4	1.498191	1.042591	0.935072	0.565841	-1.840117	-1.627426	-1.062235	0.773559	-0.497643	-0.461968
...
393	-0.856321	-0.513026	-0.479482	-0.213324	0.011586	1.621983	0.941412	0.773559	-0.497643	-0.461968
394	-0.856321	-0.925936	-1.370127	-0.993671	3.279296	1.621983	0.941412	-1.292726	-0.497643	2.164651
395	-0.856321	-0.561039	-0.531873	-0.798585	-1.440730	1.621983	0.941412	0.773559	-0.497643	-0.461968
396	-0.856321	-0.705077	-0.662850	-0.408411	1.100822	1.621983	0.941412	0.773559	-0.497643	-0.461968
397	-0.856321	-0.714680	-0.584264	-0.296088	1.391285	1.621983	0.941412	0.773559	-0.497643	-0.461968

398 rows × 10 columns

In [23]: `data.shape`

Out[23]: (398, 11)

In [26]: `x_train, x_test, y_train, y_test = train_test_split(x_s, y_s, test_size = 0.30, random_state = 1)`
`x_train.shape`

Out[26]: (278, 10)

2. Simple Linear Model

In [27]: `regression_model = LinearRegression()`
`regression_model.fit(x_train, y_train)`

```
for idx, col_name in enumerate(x_train.columns):
    print('The coefficient for {} is {}'.format(col_name, regression_model.coef_[0][idx]))

intercept = regression_model.intercept_[0]
print('The intercept is {}'.format(intercept))
```

```
The coefficient for cyl is 0.321022385691611
The coefficient for disp is 0.32483430918483897
The coefficient for hp is -0.22916950059437569
The coefficient for wt is -0.7112101905072298
The coefficient for acc is 0.014713682764191237
The coefficient for yr is 0.3755811949510748
The coefficient for car_type is 0.3814769484233099
The coefficient for origin_america is -0.07472247547584178
The coefficient for origin_asia is 0.044515252035677896
The coefficient for origin_europe is 0.04834854953945386
The intercept is 0.019284116103639764
```

2 . b Regularized Ridge Regression

```
In [28]: ridge_model = Ridge(alpha = 0.3)
ridge_model.fit(x_train, y_train)

print('Ridge model coef:{}'.format(ridge_model.coef_))
```

```
Ridge model coef:[ 0.31649043  0.31320707 -0.22876025 -0.70109447  0.01295851  0.37447352
 0.37725608 -0.07423624  0.04441039  0.04784031]
```

2. c Regularized lasso Regression

```
In [29]: lasso_model = Lasso(alpha = 0.1)
lasso_model.fit(x_train, y_train)

print('Lassomodel coef:{}'.format(lasso_model.coef_))
```

```
Lassomodel coef:[-0.          -0.          -0.01690287 -0.51890013  0.          0.28138241
 0.1278489  -0.01642647  0.          0.          ]
```

Here we noticed many coefficient are turned to 0 indicating drop of those dimensions from the model

3 . Score Comparison

```
In [32]: print(regression_model.score(x_train, y_train))
print(regression_model.score(x_test, y_test))

print('*****')
#Ridge
print(ridge_model.score(x_train, y_train))
print(ridge_model.score(x_test, y_test))

print('*****')
#Lasso
print(lasso_model.score(x_train, y_train))
print(lasso_model.score(x_test, y_test))

0.8343770256960538
0.8513421387780066
*****
0.8343617931312616
0.8518882171608507
*****
0.7938010766228453
0.8375229615977083
```

4. Model Parameter Tuning

- r^2 is not a reliable metric as it always increases with addition of more attributes even if the attributes have no influence on the predicted variable. Instead we use adjusted r^2 which removes the statistical chance that improves r^2

(adjusted $r^2 = r^2 - \text{fluke}$)

- Scikit does not provide a facility for adjusted r^2 ... so we use statsmodel, a library that gives results similar to what you obtain in R language
- This library expects the X and Y to be given in one single dataframe

```
In [33]: data_train_test = pd.concat([x_train, y_train], axis =1)
data_train_test.head()
```

	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe	mpg
350	-0.856321	-0.849116	-1.081977	-0.893172	-0.242570	1.351199	0.941412	0.773559	-0.497643	-0.461968	1.432898
59	-0.856321	-0.925936	-1.317736	-0.847061	2.879909	-1.085858	0.941412	-1.292726	-0.497643	2.164651	-0.065919
120	-0.856321	-0.695475	0.201600	-0.121101	-0.024722	-0.815074	0.941412	-1.292726	-0.497643	2.164651	-0.578335
12	1.498191	1.983643	1.197027	0.934732	-2.203196	-1.627426	-1.062235	0.773559	-0.497643	-0.461968	-1.090751
349	-0.856321	-0.983552	-0.951000	-1.165111	0.156817	1.351199	0.941412	-1.292726	2.009471	-0.461968	1.356035

```
In [34]: import statsmodels.formula.api as smf
ols1 = smf.ols(formula = 'mpg ~cyl+disp+hp+wt+acc+yr+car_type+origin_america+origin_asia', data= data_train_test).fit()
ols1.params
```

```
Out[34]: Intercept      0.019284
cyl                  0.321022
disp                 0.324834
hp                   -0.229170
wt                   -0.711210
acc                  0.014714
yr                   0.375581
car_type              0.381477
origin_america       -0.136182
origin_asia           -0.006138
dtype: float64
```

```
In [35]: print(ols1.summary())
```

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.834			
Model:	OLS	Adj. R-squared:	0.829			
Method:	Least Squares	F-statistic:	150.0			
Date:	Thu, 18 Dec 2025	Prob (F-statistic):	3.12e-99			
Time:	11:15:56	Log-Likelihood:	-146.89			
No. Observations:	278	AIC:	313.8			
Df Residuals:	268	BIC:	350.1			
Df Model:	9					
Covariance Type:	nonrobust					
coef	std err	t	P> t			
[0.025	0.975]					
Intercept	0.0193	0.025	0.765	0.445	-0.030	0.069
cyl	0.3210	0.112	2.856	0.005	0.100	0.542
disp	0.3248	0.128	2.544	0.012	0.073	0.576
hp	-0.2292	0.079	-2.915	0.004	-0.384	-0.074
wt	-0.7112	0.088	-8.118	0.000	-0.884	-0.539
acc	0.0147	0.039	0.373	0.709	-0.063	0.092
yr	0.3756	0.029	13.088	0.000	0.319	0.432
car_type	0.3815	0.067	5.728	0.000	0.250	0.513
origin_amERICA	-0.1362	0.042	-3.227	0.001	-0.219	-0.053
origin_asIA	-0.0061	0.034	-0.179	0.858	-0.074	0.061
Omnibus:	22.678	Durbin-Watson:	2.105			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36.139			
Skew:	0.513	Prob(JB):	1.42e-08			
Kurtosis:	4.438	Cond. No.	14.9			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [37]: mse = np.mean((regression_model.predict(x_test)-y_test)**2)

# root of mean_sq_error is standard deviation i.e. avg variance between predicted and actual
import math
rmse = math.sqrt(mse)
print('Root Mean Squared Error: {}'.format(rmse))
```

Root Mean Squared Error: 0.37766934254087847

In [48]: *# Is OLS a good model ? Lets check the residuals for some of these predictor.*


```
fig = plt.figure(figsize=(10,8))
sns.residplot(x= x_test['hp'], y= y_test['mpg'], color='green', lowess=True )

fig = plt.figure(figsize=(10,8))
sns.residplot(x= x_test['acc'], y= y_test['mpg'], color='green', lowess=True )
```


Out[48]: <Axes: xlabel='acc', ylabel='mpg'>


In [49]: `plt.show()`

<Figure size 1000x800 with 0 Axes>



<Figure size 1000x800 with 0 Axes>






```
In [43]: y_pred = regression_model.predict(x_test)  
plt.scatter(y_test['mpg'], y_pred)
```

```
Out[43]: <matplotlib.collections.PathCollection at 0x1fd99e72210>
```

```
In [44]: plt.show()
```



5. Inference

Both Ridge & Lasso regularization performs very well on this data, though Ridge gives a better score. The above scatter plot depicts the correlation between the actual and predicted mpg values.

This kernel is a work in progress.